

## Chapter 17

# Factoring Finite State Machines

Factoring a state machine is the process of splitting the machine into two or more simpler machines. Factoring can greatly simplify the design of a state machine by separating orthogonal aspects of the machine into separate FSMs where they can be handled independently. The separate FSMs communicate via logic signals. One FSM provides input control signals to another FSM and senses its output status signals. Such factoring, if done properly, makes the machine simpler and also makes it easier to understand and maintain — by separating issues.

In a factored FSM, the state of each sub-machine represents one dimension of a multi-dimensional state space. Collectively the states of all of the sub-machines define the state of the overall machine — a single point in this state space. The combined machine has a number of states that is equal to the product of the number of states of the individual sub-machines — the number of points in the state space.<sup>1</sup> With individual sub-machines having a few 10s of states, it is not unusual for the overall machine to have thousands to millions of states. It would be impractical to handle such a large number of states without factoring.

We have already seen one form of factoring in Section 16.3 where we developed a state machine with a data path component and a control component. In effect we factored the total state of the machine into a datapath portion and a control portion. Here we generalize this concept by showing how the control portion itself can be factored.

In this chapter we illustrate factoring by working two examples. In the first example, we start with a flat FSM and factor it into multiple simpler FSMs. In the second example we derive a factored FSM directly from the specification, without bothering with the flat FSM. Most real FSMs are designed using the latter method. A factoring is usually a natural outgrowth of the specification

---

<sup>1</sup>For most machines, not all points in the state space are reachable.




Figure 17.1: State diagram for the light flasher.

of a machine. It is rarely applied to an already flat machine.

## 17.1 A Light Flasher

Suppose you have been asked to design a light *flasher*. The flasher has a single input *in* and a single output *out*. When *in* goes high (for one cycle) it initiates the flashing sequence. During this sequence output *out*, which drives a light-emitting diode (LED) *flashes* three times. For each flash *out* goes high (LED on) for six cycles. Output *out* goes low for four cycles between flashes. After the third flash your FSM returns to the *off* state awaiting the next pulse on *in*.

A state diagram for this light flasher is shown in Figure 17.1. The FSM contains 27 states six for each of the three flashes, four for each of the two periods between flashes, and one **OFF** state. We can implement this FSM with a big **case** statement, with 27 cases. However, suppose the specifications change to require 12 cycles for each flash, 4 flashes, and 7 cycles between flashes. Because this machine is *flat* any one of these changes would require completely changing the **case** statement.

We will illustrate the process of factoring with our light flasher. This machine

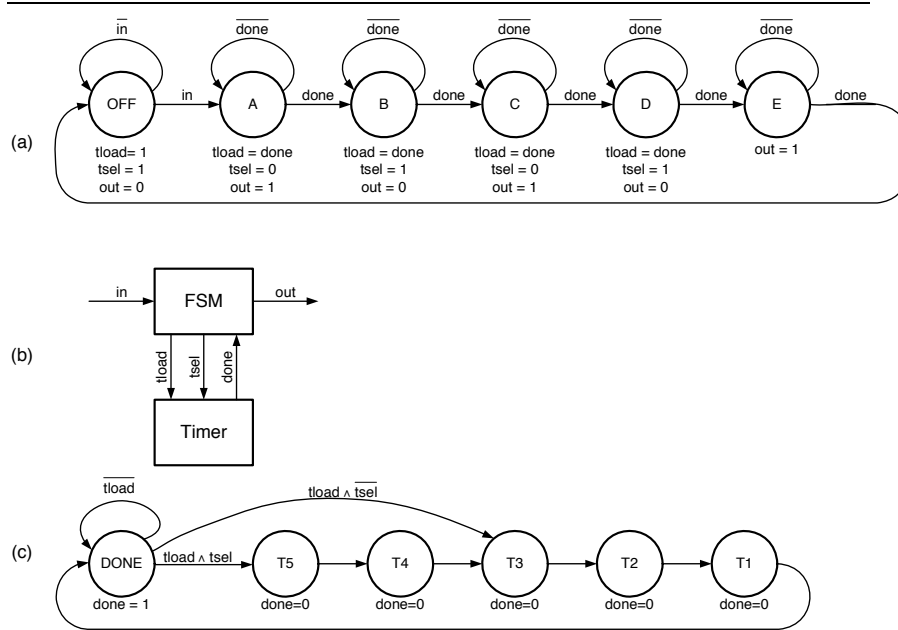


Figure 17.2: State diagram for the light flasher. With on and off timing factored out into a timer. (a) State diagram of *master* FSM. (b) Block diagram showing connections between master FSM and timer FSM. (c) State diagram of timer FSM.

can be factored in two ways. First, we can factor out the counting of on and off intervals into a timer. This allows each of the six or four state sequences for a given flash or interval to be reduced to a single state. This factoring of timing both simplifies the machine and makes it considerably simpler to change the on and off intervals. Second, we can factor out the three flashes (even though the last is slightly different). This again simplifies the machine and also allows us to change the number of flashes.

Figure 17.2 shows how the flasher is factored to separate the process of counting time intervals into a separate machine. As shown in Figure 17.2(b), the factored machine consists of a *master* FSM that accepts  $\text{in}$  and generates  $\text{out}$  and a timer FSM. The timer FSM receives two control signals from the master and returns a single status signal. The  $\text{tload}$  control signal instructs the timer to load a value into a count-down timer (i.e., to start the timer), and  $\text{tsel}$  selects the value to be loaded. When  $\text{tsel}$  is high, the timer is set to count six cycles. The counter counts four cycles when  $\text{tsel}$  is low. When the counter has completed its count down the  $\text{done}$  signal is asserted and remains asserted until the counter is reloaded. A state diagram for this timer is shown

in Figure 17.2(c).<sup>2</sup> The timer is implemented as a datapath FSM as discussed in Section 16.1.3.

A state diagram of the master FSM is shown in Figure 17.2(a). The states correspond exactly to the states of Figure 17.1 except that each sequence of repeated states (e.g., A1 to A6) is replaced by a single state (e.g., A). The machine starts in the **OFF** state. This state repeatedly loads the timer with the count-down value for the *on* sequence (**tsel** = 1). When **in** goes true, the FSM enters state A, **out** goes high, and the timer begins counting down. The master FSM stays in state A waiting for the timer to finish counting and signal **done**. It will remain in this state for six cycles. During the last cycle in state A, **done** is true, so **tload** is asserted (**tload** = **done**) and the counter is loaded with the count-down value for the *off* sequence (**tsel** = 0). Note that **tload** is asserted only during this final cycle of state A, when **done** is true. If it was asserted during each cycle of state A the timer would continually reset and never reach done. Since **done** is true in this final cycle, the machine enters state B on the next cycle and **out** goes low. The process repeats in states B through E with each state waiting for **done** before moving to the next state. Each of these states (except E) loads the counter for the next state during its last cycle by specifying **tload** = **done**.

If we compare the FSM of Figure 17.2 to the flat FSM of Figure 17.1 we realize that the state of the flat machine has been separated. The portion of the state that represents the cycle count during the current flash or space (the numeric part of the state name, or the horizontal position in Figure 17.1) is held in the counter, while the portion of state that reflects which flash or space we are currently on (the alpha part of the state name, or the vertical position in Figure 17.1) is held in the master FSM. By decomposing the state into horizontal and vertical in this manner, we are able to represent all 27 states of the original machine with just two six-state machines.

Verilog code for the master FSM of Figure 17.2(a) is shown in Figure 17.3. The **Flash** module instantiates a state register and a timer (see Figure ??) and then uses a case statement to describe the combinational next state and output logic. For each of the six states a single concatenated assignment is used to set the output signal, **out**, the two timer controls, **tload** and **tsel**, and the next state **next1**. Note that the values assigned to timer control **tload** depend on the value of the timer status **done**. This is an example of a state machine where an input **done** directly affects an output **tload** with no delay. In each state, the next state is determined with a **? :**  statement that waits for either **in** or **done** before advancing to the next state. A final assign statement resets the machine to the **OFF** state.

For completeness the verilog code for the timer is shown in Figure 17.4. The approach taken is similar to that described in Section 16.1.3.

The waveform display from a simulation of the factored flasher of Figure 17.2 is shown in Figure 17.5. The output, the fourth line from the top, shows the

---

<sup>2</sup>This state diagram only shows **tload** being active in state **DONE**. In fact, our timer can be loaded in any state. The additional edges are omitted for clarity.

---

```

// define states for flash1
`define SWIDTH 3
`define S_OFF 3'b000 // off state
`define S_A 3'b001 // first flash
`define S_B 3'b010 // first space
`define S_C 3'b011 // second flash
`define S_D 3'b100 // second space
`define S_E 3'b101 // third and final flash

// Flash - flashes out three times 6 cycles on, 4 cycles off
//      each time in is asserted.
module Flash(clk, rst, in, out) ;
    input clk, rst, in ; // in triggers start of flash sequence
    output out ; // out drives LED
    reg out ; // output
    wire ['SWIDTH-1:0] state, next ; // current state
    reg ['SWIDTH-1:0] next1 ; // next state without reset
    reg tload, tsel ; // timer inputs
    wire done ; // timer output

    // instantiate state register
    DFF #('SWIDTH) state_reg(clk, next, state) ;

    // instantiate timer
    Timer1 timer(clk, rst, tload, tsel, done) ;

    // next state and output logic
    always @(state or rst or in or done) begin
        case(state)
            `S_OFF: {out, tload, tsel, next1} =
                {1'b0, 1'b1, 1'b1, in ? `S_A : `S_OFF } ;
            `S_A: {out, tload, tsel, next1} =
                {1'b1, done, 1'b0, done ? `S_B : `S_A } ;
            `S_B: {out, tload, tsel, next1} =
                {1'b0, done, 1'b1, done ? `S_C : `S_B } ;
            `S_C: {out, tload, tsel, next1} =
                {1'b1, done, 1'b0, done ? `S_D : `S_C } ;
            `S_D: {out, tload, tsel, next1} =
                {1'b0, done, 1'b1, done ? `S_E : `S_D } ;
            `S_E: {out, tload, tsel, next1} =
                {1'b1, done, 1'b1, done ? `S_OFF : `S_E } ;
            default: {out, tload, tsel, next1} =
                {1'b1, done, 1'b1, done ? `S_OFF : `S_E } ;
        endcase
    end

    // reset to off state
    assign next = rst ? `S_OFF : next1 ;
endmodule

```

---

Figure 17.3: Verilog description of the master FSM from Figure 17.2(a).

---

```

// define time intervals
// load 5 for 6-cycle interval 5 to 0.
`define T_WIDTH 3
`define T_ON 3'd5
`define T_OFF 3'd3

// Timer 1 - reset to done state. Load time when tload is asserted
// Load with T_ON if tsel, otherwise T_OFF. If not being loaded or
// reset, timer counts down each cycle. Done is asserted and timing
// stops when counter reaches 0.
module Timer1(clk, rst, tload, tsel, done) ;
    parameter n='T_WIDTH ;
    input clk, rst, tload, tsel ;
    output done ;
    wire [n-1:0] count ;
    reg [n-1:0] next_count ;
    wire done ;

    // state register
    DFF #(n) state(clk, next_count, count) ;

    // signal done
    assign done = !(count) ;

    // next count logic
    always@(rst or tload or tsel or done or count) begin
        casex({rst, tload, tsel, done})
            4'b1xxx: next_count = 'T_WIDTH'b0 ;
            4'b011x: next_count = 'T_ON ;
            4'b010x: next_count = 'T_OFF ;
            4'b00x0: next_count = count - 1'b1 ;
            4'b00x1: next_count = count ;
            default: next_count = count ;
        endcase
    end
endmodule

```

---

Figure 17.4: Verilog description for the timer FSM used by the light flasher of Figure 17.2.

---




Figure 17.5: A waveform display showing a simulation of the factored light flasher of Figure 17.2.

desired three pulses with each pulse six clocks wide and spaces four clocks wide. The state of the master FSM is directly below the output, and the state of the timer is on the bottom line. The waveforms show how the master FSM stays in one state while the timer counts down — from 5 to 0 for flashes, and from 3 to 0 for spaces. The timer control lines (directly above the timer state) show how in states A-E (1-5) `tload` follows `done`.

We can factor our flasher further by recognizing that states A, C, and E are repeating the same function. The only difference is the number of remaining flashes. We can factor the number of remaining flashes out into a second counter as shown in Figure 17.6. Here the master FSM has only three states that determine whether the machine is off, in a flash, or in a space. The state that determines the position within a flash or space is held in a timer (just as in Figure 17.2). Finally, the state that determines the number of remaining flashes is held in a counter. Collectively the three FSMs, the master, the timer, and the counter, determine the total state of the factored machine. Each of the three sub-machines determines the state along one axis of a three-dimensional state space.

The state diagram of the master FSM for the doubly factored machine is shown in Figure 17.6(b). The machine has only three states. It starts in the **OFF** state. In the off state, both the timer and the counter are loaded. The timer is loaded with the count-down value for a flash. The counter is loaded with one less than the number of flashes required (i.e. the counter is loaded with a 3 for 4 flashes). Having input `in` go high causes a transition to the **FLASH** state. In the **FLASH** state the output `out` is true, the timer counts down, and the counter is idle. During the last cycle of the **FLASH** state the timer has reached its zero state and `tdone` is true. During this cycle the timer is reloaded with the count-down value for a space. With `tdone` true, the FSM proceeds from the **FLASH** state to the **SPACE** state if the counter is not done (`cdone` false). Otherwise, if this was the last flash (`cdone` true) the machine returns to the **OFF** state. In the **SPACE** state, the output is false, the timer counts down, and the counter is idle. In the final cycle of the **SPACE** state, `tdone` is true. This causes the counter to decrement, reducing the count of the number of remaining flashes, and the timer to reload with the count-down value for a flash.

The Verilog code for the doubly factored flasher of Figure 17.6 is shown in Figure 17.7 and the Verilog description of the counter module is shown in Figure 17.8. The next-state and output function for the master FSM again uses a concatenated assignment to set the next state, the output, and the four control signals with a single assignment. A nested `? :` statement is used to compute the next state for the **FLASH** state to test both `tdone` and `cdone`. In several states, status signal `tdone` is passed directly through to control signals `tload` and `cdec`. The counter module is nearly identical to the timer module but with slightly different control because the counter only decrements when `cdec` is asserted while the timer always decrements. With some generalization a single parameterized module could be used for both functions.

The waveforms from a simulation of the doubly factored light flasher are shown in Figure 17.9. For this simulation, the counter was initialized with a



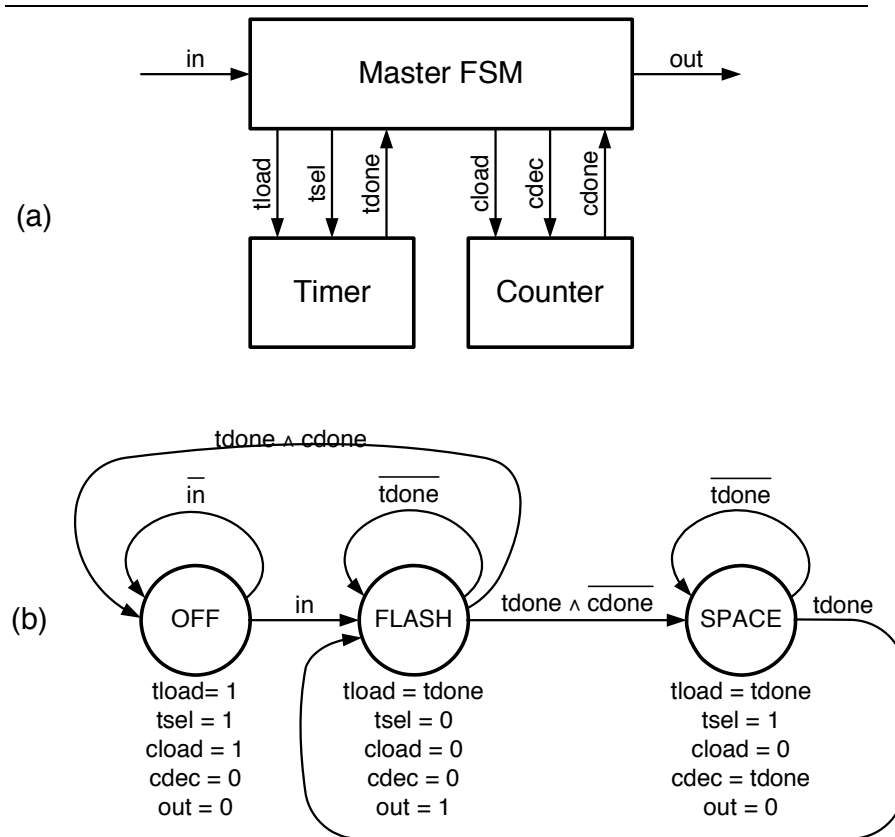


Figure 17.6: The light flasher of Figure 17.1 factored twice. The position within the current flash is held in a timer. The number of flashes remaining is held in a counter. Finally, whether we are off, in a flash, or a space is determined by the master FSM. (a) Block diagram of the twice factored machine. (b) State diagram of the master FSM.

---

```

// defines for doubly factored states
`define XWIDTH 2
`define X_OFF 2'b00
`define X_FLASH 2'b01
`define X_SPACE 2'b10

module Flash2(clk, rst, in, out) ;
    input clk, rst, in ; // in triggers start of flash sequence
    output out ; // out drives LED
    reg out ; // output
    wire ['XWIDTH-1:0] state, next ; // current state
    reg ['XWIDTH-1:0] next1 ; // next state without reset
    reg tload, tsel, cload, cdec ; // timer and countr inputs
    wire tdone, cdone ; // timer and counter outputs

    // instantiate state register
    DFF #(`XWIDTH) state_reg(clk, next, state) ;

    // instantiate timer and counter
    Timer1 timer(clk, rst, tload, tsel, tdone) ;
    Counter1 counter(clk, rst, cload, cdec, cdone) ;

    always @(state or rst or in or tdone or cdone) begin
        case(state)
            `X_OFF: {out, tload, tsel, cload, cdec, next1} =
                {1'b0, 1'b1, 1'b1, 1'b1, 1'b0,
                 in ? `X_FLASH : `X_OFF } ;
            `X_FLASH:{out, tload, tsel, cload, cdec, next1} =
                {1'b1, tdone, 1'b0, 1'b0, 1'b0,
                 tdone ? (cdone ? `X_OFF : `X_SPACE) : `X_FLASH } ;
            `X_SPACE:{out, tload, tsel, cload, cdec, next1} =
                {1'b0, tdone, 1'b1, 1'b0, tdone,
                 tdone ? `X_FLASH : `X_SPACE } ;
            default:{out, tload, tsel, cload, cdec, next1} =
                {1'b0, tdone, 1'b1, 1'b0, tdone,
                 tdone ? `X_FLASH : `X_SPACE } ;
        endcase
    end

    assign next = rst ? `X_OFF : next1 ;
endmodule

```

---

Figure 17.7: Verilog description of the master FSM from Figure 17.6.

---

```
// defines for pulse counter
// load with 3 for four pulses
`define C_WIDTH 2
`define C_COUNT 3

// Counter1 - pulse counter
//  cload - loads counter with C_COUNT
//  cdec  - decrements counter by one if not already zero
//  cdone - signals when count has reached zero

module Counter1(clk, rst, cload, cdec, cdone) ;
    parameter n=`C_WIDTH ;
    input clk, rst, cload, cdec ;
    output cdone ;
    wire [n-1:0] count ;
    reg  [n-1:0] next_count ;
    wire cdone ;

    // state register
    DFF #(n) state(clk, next_count, count) ;

    // signal done
    assign cdone = !(count) ;

    // next count logic
    always@(rst or cload or cdec or cdone or count) begin
        casex({rst, cload, cdec, cdone})
            4'b1xxx: next_count = `C_WIDTH'b0 ;
            4'b01xx: next_count = `C_COUNT ;
            4'b0010: next_count = count - 1'b1 ;
            4'b00x1: next_count = count ;
            default: next_count = count ;
        endcase
    end
endmodule
```

---

Figure 17.8: Verilog description of the counter from Figure 17.6.

---




Figure 17.9: A waveform display showing a simulation of the doubly factored light flasher of Figure 17.6.

3 so that the master FSM produces 4 flashes. The different time scales of the three state variables are clearly visible. The counter (bottom line) moves most slowly, counting down from 3 to 0 over the four-flash sequence. It decrements after the last cycle of each space. At each point in time, it represents the number of flashes left to do after the current flash is complete. The master FSM state moves the next most slowly. After starting on the 0 (OFF) state, it alternates between the 1 (FLASH) state and the 2 (SPACE) state until the four flashes are complete. Finally, the timer state moves most rapidly counting down from 5 to 0 for flashes, or 3 to 0 for spaces.

## 17.2 Traffic Light Controller

As a second example of factoring, we consider a more sophisticated version of the traffic-light controller we introduced in Section 14.3. This machine has two inputs `car_ew` and `car_lt` that indicate that cars are waiting on the east-west road (`ew`) and that cars are waiting in a left-turn lane (`lt`). The machine has nine output lines that drive three sets of three lights each one set each for the north-south road, the east-west road, and the left-turn lane (from the north-south road). Each set of lights consists of a red light, a yellow light, and a green light.

Normally the light will be green for the north-south road. However, if a car is detected on the east-west road or the left-turn lane, we wish to switch the lights so that the east-west or left-turn light is green (with priority going to the left-turn lane). Once the lights are switched to east-west or left-turn, we leave them switched until either no more cars are detected in that direction, or until a timer expires. The lights then return to green in the north-south direction.

Each time we switch the lights we switch the lights in the active direction from green to yellow. Then, after a time interval, we switch them to red. Then, after a second time interval, we switch them to green. The lights are not allowed to change again until they have been green for a third time interval.

Given this specification we decide to factor the finite-state machine that implements this traffic-light controller into five modules as shown in Figure 17.10. A master FSM accepts the inputs and decides which direction should be green `dir`. To time when it is time to force the lights back to north-south, it uses a timer, `Timer1`. It also receives a signal `ok` back from the combiner that indicates that the sequence from the last direction change is complete and the direction is allowed to change again.

The combiner module maintains a current direction state, and combines this current direction state with the `light` signal from the light FSM to generate the 9 light outputs `lights`. The combiner also receives direction requests from the master FSM on the `dir` signal and sequences the light FSM in response to these requests. When a new direction request occurs, the combiner deasserts `on` (sets it low) to ask the light FSM to switch the lights to red. Once the lights are red, the current direction is set equal to the direction request and the `on` signal is asserted to request that the lights be sequenced to green. Only after




Figure 17.10: Block diagram of a factored traffic-light controller.

the **done** signal from the light FSM is asserted, signaling that the sequence is complete and the lights have been green for the required time interval is the **ok** signal asserted to allow another direction change.

The modules here illustrate two types of relationships. The master FSM and combiner modules form a *pipeline*. Requests flow down this pipeline from left to right. A request is input to the Master FSM in the form of a transition on the **car\_ew** and **car\_lt** inputs. The Master processes this request and in turn issues a request to the combiner on the **dir** lines. The combiner in turn processes the request and outputs the appropriate sequence on the **lights** output in response. We will discuss pipelines in more depth in Chapter 21.

The **ok** signal here is an example of a *flow control* signal. It provides back-pressure on the master FSM to prevent it from getting ahead of the combiner and light FSMs. The master FSM makes a request and it is not allowed to make another request until the **ok** signal indicates that the rest of the circuit is finished processing the first request.

The other relationships in Figure 17.10 are *master-slave* relationships. The Master FSM acts as a master to Timer1, giving it commands, and Timer1 acts as a slave, receiving the commands and carrying them out. In a similar manner the combiner is the master of the Light FSM which in turn is the master of Timer2.

The light FSM sequences the traffic lights from green to red and then back to green. It receives requests from the combiner on the **on** signal and responds to these requests with the **done** signal. It also generates the 3-bit **light** signal which indicates which light in the current direction should be illuminated. When




Figure 17.11: State diagram for the master FSM of Figure 17.10.

the `on` signal is set high, it requests that the light signal switch the `light` signal to green. When this is completed, and the minimum green time interval has elapsed, the light FSM responds by asserting `done`. When the `on` signal is set low, it requests that the light FSM sequence the `light` signal to red — via a yellow state and observing the required time intervals. When this is completed `done` is set low. The light FSM uses its own timer (Timer2) to count out the time intervals required for light sequencing.

For the interface between the light FSM and the combiner, the `done` signal provides flow control. The combiner can only toggle `on` high when `done` is low and can only toggle `on` low when `done` is high. After toggling `on`, it must wait for `done` to switch to the same state as `on` before switching `on` again.

A state diagram for the master FSM is shown in Figure 17.11. The machine starts in the NS state. In this state Timer1 is loaded so it can count down in the LT and EW states, and the requested direction is NS. State NS is exited when the `ok` signal indicates that a new direction can be requested and when one of the `car` signals indicates that there is a car waiting in another direction. In the EW and LT states, the new direction is requested and the state is exited with `ok` is true and either there is no longer a car in that direction or the timer has expired, as signaled by `tdone`.

The light FSM is a simple light sequencer similar to our light flasher of Section 17.1. A state diagram for the light FSM is shown in Figure 17.12. Like the light flasher, during the last cycle in each state the timer is loaded with the timeout for the next state. The machine starts in the RED state. It transitions to GREEN when the timer is done and the `on` signal from the combiner indicates




Figure 17.12: State diagram for the light FSM of Figure 17.10.

---

that a green light is requested. The **done** signal is asserted after the timer in the **GREEN** state has completed its count down. The transition to **YELLOW** is triggered by the timer being complete and the **on** signal being low. The transition from **YELLOW** to **RED** occurs on the timer alone. The **done** signal is held high during the **YELLOW** state. It is allowed to go low, signaling that the transition to **RED** is complete, only after the timer has completed its countdown in the **RED** state.

A Verilog description of the master FSM for the factored traffic-light controller is shown in Figure 17.13. This module instantiates a timer and then uses a **case** statement to realize a three-state FSM that exactly follows the state diagram of Figure 17.11. Verilog select, “? :”, statements are used for the next-state logic. A triply-nested select statement is used in state **M\_NS** to test in turn signals **ok**, **car\_lt**, and **car\_ew**.

The master FSM does not instantiate the combiner. Both the master and combiner are instantiated as *peer* modules at the top level.

A Verilog description of the combiner module is shown in Figure 17.14. This module accepts a **dir** input from the master FSM, responds to the master FSM with flow-control signal **ok**, and generates the **lights** output. The key piece of state in the combiner module is the current direction register **cur\_dir**. This holds the direction that the light FSM is currently sequencing. It is updated with the requested direction, **dir**, when **on** and **done** are both low. This occurs when the combiner has requested that the lights be sequence to red (**on** low), and the light FSM has completed this requested action (**done** low).

The **on** command, which requests the light FSM to sequence the lights to green, is asserted any time the current direction and the requested direction match. When the master FSM requests a new direction, this causes **on** to go low - requesting the light FSM to sequence the lights red in preparation for the new direction.

The **ok** response to the master FSM is asserted when **on** and **done** are both true. This occurs when the light FSM has completed sequencing the requested direction to green.

The **lights** output is computed by a **case** statement with the current direction as the case variable. This case statement inserts the **light** output from



---

```

// master FSM states
// these also serve as direction values
`define MWIDTH 2
`define M_NS    2'b00
`define M_EW    2'b01
`define M_LT    2'b10

//Master FSM
// car_ew - car waiting on east-west road
// car_lt - car waiting in left-turn lane
// ok      - signal that it is ok to request a new direction
// dir     - output signaling new requested direction

module TLC_Master(clk, rst, car_ew, car_lt, ok, dir) ;
    input clk, rst, car_ew, car_lt, ok ;
    output [1:0] dir ;

    wire [MWIDTH-1:0] state, next ; // current state and next state
    reg  [MWIDTH-1:0] next1 ;        // next state without reset
    reg  tload ;                     // timer load
    reg  [1:0] dir ;                  // direction output
    wire tdone ;                     // timer completion

    // instantiate state register
    DFF #('MWIDTH) state_reg(clk, next, state) ;

    // instantiate timer
    Timer #('TWIDTH) timer(clk, rst, tload, 'T_EXP, tdone) ;

    always @(state or rst or car_ew or car_lt or ok or tdone) begin
        case(state)
            'M_NS: {dir, tload, next1} =
                {'M_NS, 1'b1, ok ? (car_lt ? 'M_LT
                                     : (car_ew ? 'M_EW : 'M_NS))
                 : 'M_NS} ;
            'M_EW: {dir, tload, next1} =
                {'M_EW, 1'b0, (ok & (!car_ew | tdone)) ? 'M_NS : 'M_EW} ;
            'M_LT: {dir, tload, next1} =
                {'M_LT, 1'b0, (ok & (!car_ew | tdone)) ? 'M_NS : 'M_LT} ;
            default: {dir, tload, next1} =
                {'M_NS, 1'b0, 'M_NS} ;
        endcase
    end
    assign next = rst ? 'M_NS : next1 ;
endmodule

```

---

Figure 17.13: Verilog description of the master FSM for the traffic-light controller.

---

```

//-----
// Combiner -
//  dir - direction request from master FSM
//  ok  - acknowledge to master FSM
//  lights - 9-bits to control traffic lights {NS,EW,LT}
//-----
module TLC_Combiner(clk, rst, dir, ok , lights) ;
    input clk, rst ;
    input [1:0] dir ;
    output ok ;
    output [8:0] lights ;
    wire done ;
    wire [2:0] light ;
    reg [8:0] lights ;
    wire [1:0] cur_dir ;

    // request green from light FSM until direction changes
    wire on = (cur_dir == dir) ;

    // update direction when light FSM has made lights red
    wire [1:0] next_dir = rst ? 2'b0 : ((!on & !done) ? dir : cur_dir) ;

    // ok to take another change when light FSM is done
    wire ok = on & done ;

    // current direction register
    DFF #(2) dir_reg(clk, next_dir, cur_dir) ;

    // combine cur_dir and light to get lights
    always @(cur_dir or light) begin
        case(cur_dir)
            'M_NS: lights = {light, 'RED, 'RED} ;
            'M_EW: lights = {'RED, light, 'RED} ;
            'M_LT: lights = {'RED, 'RED, light} ;
            default: lights = {'RED, 'RED, 'RED} ;
        endcase
    end

    // light FSM
    TLC_Light lt(clk, rst, on, done, light) ;
endmodule

```

Figure 17.14: Verilog description of the combiner for the traffic-light controller.

the light FSM into the position corresponding to the current direction and sets the other positions to be red.

A Verilog description of the light FSM is shown in Figure 17.15. The module instantiates a timer and then uses a **case** statement to implement a FSM with state transitions as shown in Figure 17.12. Verilog select statements are used for the next-state logic.

Waveforms from a simulation of the factored traffic-light controller are shown in Figure 17.16. The machine is initially reset with the Master FSM in the NS state (00). Output **dir** is also NS (00). The **ok** line is initially low because the light FSM has not yet finished its sequencing to make the lights green in the NS direction. The lights are initially all red (444) but change after one cycle to be green in the north-south direction (144).

The light FSM is initialized to the RED (00) state and advances to the green state because **on** and **tdone** are both asserted. In the green state it starts a timer and waits until **tdone** is again asserted before signaling **done** to the combiner which in turn causes the combiner to signal **ok** to the Master FSM.

Because **car\_ew** is asserted, once the **ok** signal goes high, the Master FSM requests a change in direction to east-west by setting **dir** to 01. The combiner responds by setting **on** low to request the light FSM to sequence the current direction's light to red. The light FSM in turn responds by transitioning to the YELLOW state (10) which causes **light** to go to yellow (2) and **lights** to go to 244 - yellow in the north south direction. When the light timer completes its countdown, the light machine enters the RED state (00), **light** becomes 4 (RED), and **lights** becomes 444 - red in all directions. Once the light timer counts down the minimum time in the all-red state, the light FSM sets **done** low signaling that it has completed the transition.

With **on** and **done** both low, the combiner updates **cur\_dir** to east-west (01) and sets **on** high to request the light FSM to sequence the lights green in the east-west direction. When the light FSM completes this action - including timing the green state - it sets **done** true. This in turn causes the combiner to set **ok** true, signaling the Master FSM that it is ready to accept a new direction.

When **ok** is asserted the second time, the Mater FSM requests the north-south direction again (**dir** = 00). This decision is driven by the **car\_ew** line being low and the master timer being done. This new direction causes **on** to go low, sequencing the lights to the all red state. Then, after **cur\_dir** is updated, setting **on** high to sequence them back to the green state. When this is all complete, **ok** is asserted again.

On this third assertion of **ok**, signal **car\_lt** is true, so a left-turn is requested (**dir** = 10) and the lights are sequenced to red and back to green again. When the sequencing is completed, **ok** is asserted for a fourth time. This time **car\_lt** is still asserted, but the master timer is done, so a north-south direction is again requested.

```

//-----
// Light FSM
//-----
module TLC_Light(clk, rst, on, done, light) ;
    input clk, rst, on ;
    output done ;
    output [2:0] light ;
    reg [2:0] light ;
    reg done ;
    wire ['LWIDTH-1:0] state, next ; // current state, next state
    reg ['LWIDTH-1:0] next1 ; // next state w/o reset
    reg tload ;
    reg ['TWIDTH-1:0] tin ;
    wire tdone ;

    // instantiate state register
    DFF #('LWIDTH) state_reg(clk, next, state) ;


    // instantiate timer
    Timer timer(clk, rst, tload, tin, tdone) ;

    always @(state or rst or on or tdone) begin
        case(state)
            'L_RED: {tload, tin, light, done, next1} =
                {tdone & on, 'T_GREEN, 'RED, !tdone,
                 (tdone & on) ? 'L_GREEN : 'L_RED} ;
            'L_GREEN: {tload, tin, light, done, next1} =
                {tdone & !on, 'T_YELLOW, 'GREEN, tdone,
                 (tdone & !on) ? 'L_YELLOW : 'L_GREEN} ;
            'L_YELLOW: {tload, tin, light, done, next1} =
                {tdone, 'T_RED, 'YELLOW, 1'b1, tdone ? 'L_RED : 'L_YELLOW} ;
            default: {tload, tin, light, done, next1} =
                {tdone, 'T_RED, 'YELLOW, 1'b1, tdone ? 'L_RED : 'L_YELLOW} ;
        endcase
    end

    assign next = rst ? 'L_RED : next1 ;
endmodule

```

Figure 17.15: Verilog description of the light FSM for the traffic-light controller.



### 17.3 Exercises

- 17-1 *Flasher*. Modify the flasher FSM to flash an SOS sequence - three short flashes (one clock each) followed by three long flashes (four clocks each) followed by three short flashes again.
- 17-2 *Traffic Light Controller*. Modify the traffic light controller so that north-south and east-west have equal priority. Add an additional input, `car_ns` that indicates when a car is waiting in the north-south direction. In either the NS or EW states, change to the other if there is a car waiting in the other direction **and** the master timer is done.
- 17-3 *Traffic Light Controller*. Modify the traffic light controller so that a switch from a red light to a green light is preceded by both red and yellow being on for three clocks.

## Chapter 18

# Microcode

Realizing the next-state and output logic of a finite-state machine using a memory array gives a flexible way of realizing a FSM. The function of the FSM can be altered by changing the contents of the memory. We refer to the contents of the memory array as *microcode* and a machine realized in this manner is called a *microcoded* FSM. Each word of the memory array determines the behavior of the machine for a particular state and input combination and is referred to as a *microinstruction*.

We can reduce the required size of a microcode memory by augmenting the memory with special logic to compute the next state, and by selectively updating infrequently changing outputs. A single microinstruction can be provided for each state, rather than one for each state  $\times$  input combination, by providing an instruction sequencer and a *branch* microinstruction to cause changes in control flow. Bits of the microinstruction can be shared by different functions by defining different microinstruction types for control, output, and other functions.

Microcode was originated by Maurice Wilkes at Cambridge University in 1951 to implement the control logic for the EDSAC computer [?]. It has been widely used since that time in many different types of digital systems.

### 18.1 A Simple Microcoded FSM

Figure 18.1 shows a block diagram of a simple microcoded FSM. A memory array holds the next-state and output functions. Each word of the array holds the next state and output for a particular combination of input and next state. The array is addressed by the concatenation of the current state and the inputs. A pair of registers holds the current state and current output.

In practice the memory could be realized as a RAM or EEPROM allowing software to reprogram the microcode. Alternatively the memory could be a ROM. With a ROM, a new mask set is required to reprogram the microcode. However, this is still advantageous as changing the program of the ROM does not otherwise alter the layout of the chip. Some ROM designs even allow the




Figure 18.1: Block diagram of a simple microcoded FSM.

program to be changed by altering only a single metal-level mask — reducing the cost of the change. Some designs take a hybrid approach putting most of the microcode into ROM (to reduce cost) but keeping a small portion of microcode in RAM. A method is provided to redirect an arbitrary state sequence into the RAM portion of the microcode to allow any state to be *patched* using the RAM.

A verilog description of this FSM is shown in Figure 18.2. It follows the schematic exactly with the addition of some logic to reset the state when `rst` is asserted. The module `ROM` is a read-only-memory that takes an address `{state, in}` and returns a microinstruction `uinst`. The microinstruction is then split into its next-state and output components. The FSM of Figures 18.1 and 18.2 is very simple because it has no function until the ROM is programmed.

To see how to program the ROM to realize a finite state machine, consider our simple traffic-light controller from Section 14.3. The state diagram of this controller is repeated in Figure 18.3. To fill our microcode ROM, we simply write down the next state and output for each current state/input combination as shown in Table 18.1. Consider the first line of the table. Address 0000 corresponds to state GNS (green north-south) with input `car_ew` = 0. For this state, the output is 100001 (green north-south, red east-west) and the next state is GNS (000). Thus, the contents of ROM location 0000 is the 000100001, the concatenation of the next state 000 with the output. For the second line of the table, address 0001 corresponds to GNS with `car_ew` = 1. The output here is the same as for the first line, but the next state is now YNS (001) hence the contents of this ROM location is 001100001. The remaining rows of the table are derived in a similar manner. The ROM itself is loaded with the contents of the column labeled Data.

The results of simulating the microcoded FSM of Figure 18.2 using the ROM contents from Table 18.1 are shown in Figure 18.4. The output, state, and microcode ROM address, and microcode ROM data (microinstruction) (the bottom four signals) are displayed in octal (base 8). The system initializes to state 0 (GNS) with output 41 (green (4) north-south, red (1) east-west). Then



---

```

module ucode1(clk,rst,in,out) ;
  parameter n = 1 ; // input width
  parameter m = 6 ; // output width
  parameter k = 3 ; // bits of state

  input  clk, rst ;
  input  [n-1:0] in ;
  output [m-1:0] out ;

  wire  [k-1:0] next, state ;
  wire [k+m-1:0] uinst ;

  DFF #(k) state_reg(clk, next, state) ; // state register
  DFF #(m) out_reg(clk, uinst[m-1:0], out) ; // output register
  ROM #(n+k,m+k) uc({state, in}, uinst) ; // microcode store
  assign next = rst ? {k{1'b0}} : uinst[m+k-1:m] ; // reset state
endmodule

```

---

Figure 18.2: Verilog description of a simple microcoded FSM




Figure 18.3: State diagram of simple traffic-light controller.

Address	State	car_ew	Next State	Output	Data
0000	GNS (000)	0	GNS (000)	100001	000100001
0001	GNS (000)	1	YNS (001)	100001	001100001
0010	YNS (001)	0	GEW (010)	010001	010010001
0011	YNS (001)	1	GEW (010)	010001	010010001
0100	GEW (010)	0	YEW (011)	001100	011001100
0101	GEW (010)	1	YEW (011)	001100	011001100
0110	YEW (011)	0	GNS (000)	001010	000001010
0111	YEW (011)	1	GNS (000)	001010	000001010

Table 18.1: State table for simple microcoded traffic-light controller

the input line (`car_ew`) goes high switching the ROM address from 00 to 01. This causes the microinstruction to switch from 041 to 141 which selects a next state of 1 (YNS) on the next clock. The machine then proceeds through states 2 (GEW) and 3 (YEW) before returning to state 0.

The beauty of microcode is that we can change the function of our FSM by changing only the contents of the ROM. Suppose for example that we would like to modify the FSM so that:

1. The light stays green in the east-west direction as long as `car_ew` is true.
2. The light stays green in the north south direction for a minimum of three cycles (states GNS1, GNS2, and GNS3).
3. After a yellow light, the lights should go red in both directions for one cycle before turning the new light green.

Table 18.2 shows the state table that implements these modifications. We split state GNS into three states and add two new states (RNS and REW). Note that the GEW state now tests `car_ew` and stays in GEW as long as it is true. The results of simulating our FSM of Figure 18.2 with this new microcode is shown in the waveform display of Figure 18.5.

## 18.2 Instruction Sequencing

A microcoded FSM can be made considerably more efficient by using a *sequencer* to generate the address of the next instruction. Performing instruction sequencing has two big advantages. First, for microinstructions that simply proceed to the next instruction, this instruction address can be generated with a counter, eliminating the need to store the address in the microcode memory. Second, by using logic to select or combine the different inputs, the microcode store can store a single microinstruction for each state, rather than having to store separate (and nearly identical) instructions for each possible combination of inputs.

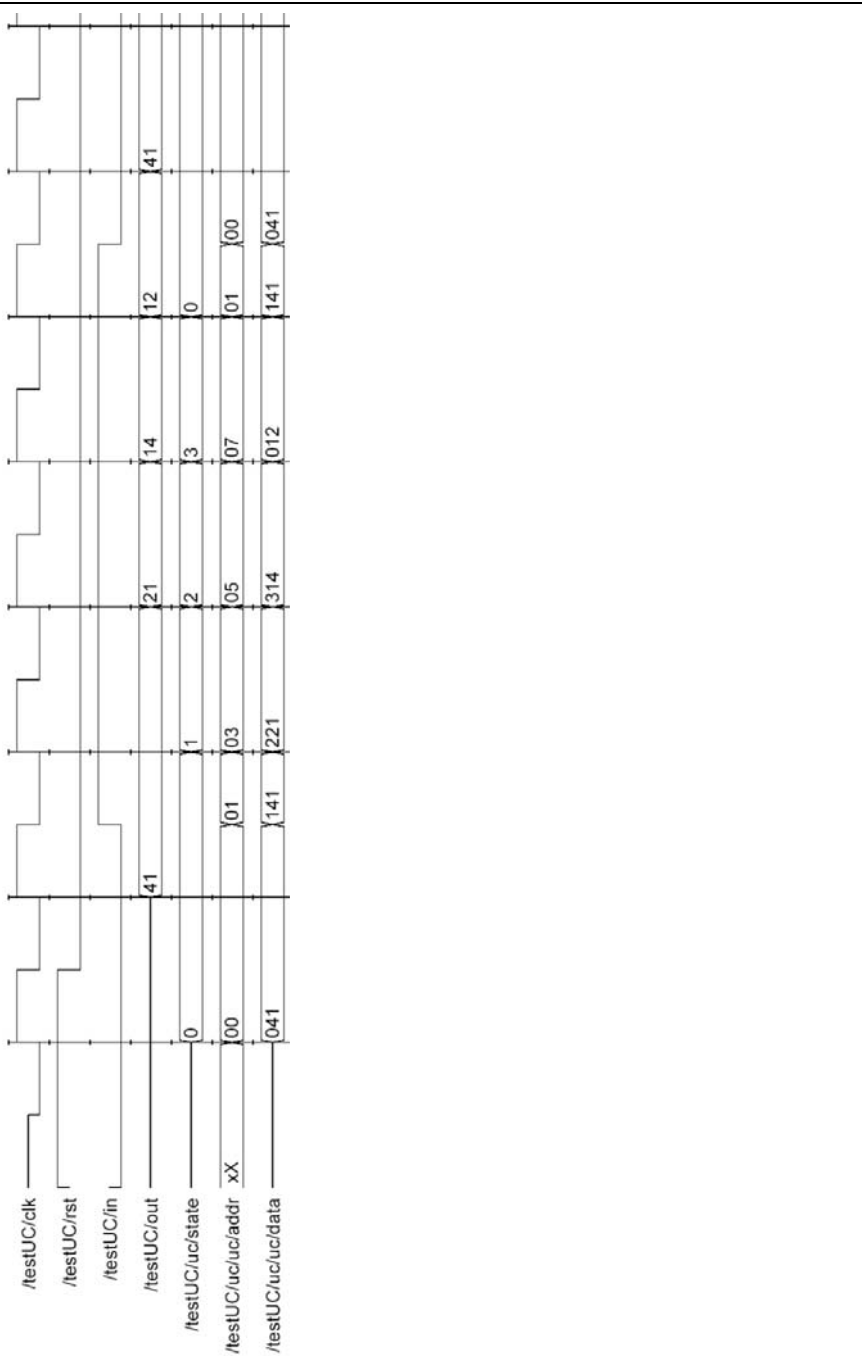


Figure 18.4: A waveform display showing a simulation of the microcoded FSM of Figure 18.2 using the microcode from Table 18.1.




Figure 18.5: A waveform display showing a simulation of the microcoded FSM of Figure 18.2 using the microcode from Table 18.2.

Address	State	car_ew	Next State	Output	Data
0000	GNS1(000)	0	GNS2(001)	100001	001100001
0001	GNS1(000)	1	GNS2(001)	100001	001100001
0010	GNS2(001)	0	GNS3(010)	100001	010100001
0011	GNS2(001)	1	GNS3(010)	100001	010100001
0100	GNS3(010)	0	GNS3(010)	100001	010100001
0101	GNS3(010)	1	YNS (011)	100001	011100001
0110	YNS (011)	0	RNS (100)	010001	100010001
0111	YNS (011)	1	RNS (100)	010001	100010001
1000	RNS (100)	0	GEW (101)	001001	101001001
1001	RNS (100)	1	GEW (101)	001001	101001001
1010	GEW (101)	0	YEW (110)	001100	110001100
1011	GEW (101)	1	GEW (101)	001100	101001100
1100	YEW (110)	0	REW (111)	001010	111001010
1101	YEW (110)	1	REW (111)	001010	111001010
1110	REW (111)	0	GNS (000)	001001	000001001
1111	REW (111)	1	GNS (000)	001001	000001001

Table 18.2: State table for simple microcoded traffic-light controller

A review of the microcode from Table 18.2 shows the redundancy a sequencer can eliminate. Each instruction includes an explicit next-state field while all instructions either select themselves or the next instruction in sequence for the next state. Also, all instructions are duplicated for both input states, with only two having a small difference in the next state field. This overhead would be even higher if there were more than one input signal.

Adding a sequencer to our microcoded FSM is the first step from an FSM (where the next state is determined by a logic function) to a stored-program computer, where the next instruction is determined by interpreting the current instruction. With a sequencer, like a stored program computer, our microcoded machine *executes* microinstructions in sequence until a branch instruction redirects execution to a new address.

Figure 18.6 shows a microcoded FSM that uses an instruction sequencer. The state register here is replaced with a microprogram counter ( $\mu$ PC or uPC) register. At any point in time this register represents the current state by selecting the current microinstruction. With this design we have reduced the number of microinstructions in the microcode memory from  $2^{s+i}$  to  $2^s$ , that is from  $2^i$  per state to 1 per state. The cost of this reduction is increasing the width of each microinstruction from  $s + o$  bits to  $s + o + b$  bits. Each microinstruction consists of three fields as shown in Figure 18.7: an  $o$ -bit field that specifies the current output, an  $s$ -bit field that specifies the address to branch to (the branch target), and a  $b$ -bit field that specifies a branch instruction.

With the instruction sequencer, inputs are tested by the branch logic as directed by a branch instruction from the current microinstruction. As a result



Encoding	OpCode	Description
000	NOP	Never branch, always proceed to uPC+1
001	B0	Branch on input 0. If <code>in[0]</code> branch to <code>br_upc</code> otherwise continue to uPC+1
010	B1	Branch on input 1.
011	BA	Branch any. Branch if either input is true.
100	BR	Always branch. Select <code>br_upc</code> as the next uPC regardless of the inputs.
101	BN0	Branch on not input 0. If <code>in[0]</code> is false, branch, otherwise continue to uPC+1.
110	BN1	Branch on not input 1.
111	BNA	Branch only if inputs 0 and 1 are both false.

Table 18.3: Branch instruction encodings.

of this test, the sequencer either branches (by selecting the branch target field as the next uPC) or doesn't (by selecting uPC+1 as the next UPC).

Consider an example with a two-bit input field. We can define a three-bit branch instruction `brinst` as follows.

```
branch = (brinst[0] & in[0] | brinst[1] & in[1]) ^ brinst[2] ;
```

Branch instruction bits 0 and 1 select whether we test input bits 0 or 1 (or either). Branch instruction bit 2 controls the polarity of the test. If `brinst[2]` is low, we branch if the selected bit(s) is high. Otherwise we branch if the selected bit(s) is low. Using this encoding of the branch instruction, we can perform the branches shown in Table 18.3.

Other encodings of the branch instruction are possible. A common  $n$ -bit encoding uses  $n - 1$ -bits to select one of  $2^{n-1}$  inputs to test and the remaining bit to select whether to branch on the selected input high or low. One of the inputs is set always high to allow the NOP and BR instructions to be created. For this encoding the branch signal would be:

```
branch = brinst[n-1] ^ in[brinst[n-2:0]] ;
```

The branch instructions created by this alternate encoding (for a 3-bit `brinst` and three inputs are listed in Table 18.4. To provide the NOP and BR instructions, we use a constant 1 for the fourth input. Here each branch instruction tests exactly one input, while in the encoding of Table 18.3 the instructions may test zero, one, or two inputs. Many other possible encodings beyond to the two presented here are possible.

A Verilog description of the microcoded FSM with an instruction sequencer is shown in Figure 18.8. The Verilog follows the block diagram of Figure 18.6 closely. One assign statement calculates signal `branch` which is true if the sequencer is to branch on the next cycle. A second assign statement then calculates the next micro program counter (`nupc`) based on `branch` and `rst`.

Encoding	OpCode	Description
000	B0	Branch on input 0.
001	B1	Branch on input 1.
010	B2	Branch on input 2.
011	BR	Always branch. (Input 3 is the constant “1”).
100	BN0	Branch on not input 0.
101	BN1	Branch on not input 1.
110	BN2	Branch on not input 2.
111	NOP	Never branch.

Table 18.4: Alternate branch instruction encodings.

---

```

module ucode2(clk,rst,in,out) ;
    parameter n = 2 ; // input width
    parameter m = 9 ; // output width
    parameter k = 4 ; // bits of state
    parameter j = 3 ; // bits of instruction

    input  clk, rst ;
    input  [n-1:0] in ;
    output [m-1:0] out ;

    wire  [k-1:0] nupc, upc ; // microprogram counter
    wire [j+k+m-1:0] uinst ; // microinstruction word

    // split off fields of microinstruction
    wire [m-1:0] nxt_out ; // = uinst[m-1:0] ;
    wire [k-1:0] br_upc  ; // = uinst[m+k-1:m] ;
    wire [j-1:0] brinst  ; // = uinst[m+j+k-1:m+k] ;
    assign {brinst, br_upc, nxt_out} = uinst ;

    DFF #(k) upc_reg(clk, nupc, upc) ; // microprogram counter
    DFF #(m) out_reg(clk, nxt_out, out) ; // output register
    ROM #(k,m+k+j) uc(upc, uinst) ; // microcode store

    // branch instruction decode
    wire branch = (brinst[0] & in[0] | brinst[1] & in[1]) ^ brinst[2] ;

    // sequencer
    assign nupc = rst ? {k{1'b0}} : branch ? br_upc : upc + 1'b1 ;
endmodule

```

---

Figure 18.8: Verilog description of a microcoded FSM with an instruction sequencer



Address	State	Br Inst	Target	NS LT EW	Data
0000	NS1	BLT (001)	LT1(0101)	100001001	0010101100001001
0001	NS2	BNEW (110)	NS1(0000)	100001001	1100000100001001
0010	EW1	NOP (000)		010001001	0000000010001001
0011	EW2	BEW (010)	EW2(0011)	001001100	0100011001001100
0100	EW3	BR (100)	NS1(0000)	001001010	1000000001001010
0101	LT1	NOP (000)		010001001	0000000010001001
0110	LT2	BLT (001)	LT2(0110)	001100001	0010110001100001
0111	LT3	BR (100)	NS1(0000)	001010001	1000000001010001

Table 18.5: Microcode for traffic light controller with sequencer of Figure 18.6.

To make the code more readable we use an assign statement to split out the three fields of the microinstruction into the output (**next\_out**), the branch target (**br\_upc**), and the branch instruction (**brinst**). Using these mnemonic names, rather than indexing fields of **uinst** makes the remainder of the code easier to follow.

Consider a slightly more involved version of our traffic-light controller that includes a left-turn signal as well as north-south and east-west signals. Table 18.5 shows the microcode. Here input 0 is **car\_lt** and input 1 is **car\_ew** so we rename our branches BLT (branch if **car\_lt**), BNEW (branch if not **car\_ew**) and so on.

The microcode of Table 18.5 starts in state NS1 where the light is green in the north-south direction. In this state the left-turn sensor is checked with a BLT to LT1. If **car\_lt** is true, control transfers to LT1. Otherwise the uPC proceeds to the next state, NS2. In NS2 the microcode branches back to NS1 if **car\_ew** is false (BNEW NS1). Otherwise control falls through to EW1 where the north-south light goes yellow. EW1 is always followed by EW2 where the east-west light is green. A BEW EW2 keeps the uPC in state EW2 as long as **car\_ew** is true. When **car\_ew** goes false, the uPC proceeds to state EW3 where the east-west light is yellow and a BR NS1 transfers control back to NS1. The left turn sequence (LT1, LT2, LT3) operates in a similar manner.

Waveforms from a simulation of the microcoded sequencer of Figure 18.8 running the microcode of Table 18.5 is shown in Figure 18.9. The fifth row from the top shows the microprogram counter **upc**. The machine is reset to **upc = 0** (NS1) advances to 1 (NS2), then back to 0 (NS1) before branching to 5 (LT1). It proceeds from 5 (LT1) to 6 (LT2) and remains in 6 until **car\_lt** goes low. At that point it advances to 7 (LT3) and back to 0 (NS1). At this point **car\_ew=1**, and the sequence followed is 0, 1, 2, 3 (NS1, NS2, EW1, EW2). The machine stays in 3 (EW2) until **car\_ew** goes low and then proceeds to 4 (EW3) and back to 0 (NS1). The machine cycles between NS1 and NS2 a few times until both **car\_ew** and **car\_lt** go high at the same time. As the machine is in NS1 when this happens, **car\_lt** is checked first and the uPC is directed to LT1.

Because the microcoded FSM of Figure 18.6 can only branch one way in each microinstruction, it takes two states (NS1 and NS2) to perform a three-




Figure 18.9: A waveform display showing a simulation of the microcoded FSM of Figure 18.2 using the microcode from Table 18.2.

Address	State	Br Inst	Target	NS LT EW	Data
0000	NS1	BNA (111)	NS1(0000)	100001001	1110000100001001
0001	NS2	BLT (001)	LT1(0100)	010001001	0010100010001001
0010	EW1	BEW (010)	EW1(0010)	001001100	0100010001001100
0011	EW2	BR (100)	NS1(0000)	001001010	1000000001001010
0100	LT1	BLT (001)	LT1(0100)	001100001	0010100001100001
0101	LT2	BR (100)	NS1(0000)	001010001	1000000001010001

Table 18.6: Alternate microcode for traffic light controller with sequencer of Figure 18.6.

way branch between staying with north-south, going to east-west, or going to left-turn. This results in two states with the lights green in the north-south direction (NS1 and NS2) and two states with the lights yellow in the north-south direction (EW1 and LT1). The real solution to this problem is to support a multi-way branch (which we will discuss below). However we can partially solve the problem in software by using the alternate microcode shown in Table 18.6.

In the alternate microcode of Table 18.6, the uPC stays in state NS1 as long as `car_ew` and `car_lt` are both false by using a BNA NS1 (branch on not any inputs to NS1). NS1 is now the only state with the lights green in the north-south direction. If any inputs are true, the uPC proceeds to state NS2 which is the single state with the lights yellow in the north-south direction. State NS2 tests the `car_lt` input and branches to state LT1 if true (BLT LT1). If `car_lt` is false, the uPC proceeds to EW1. The remainder of the machine is similar to that of Table 18.5 except that the EW and LT states have been renumbered.

Simulation waveforms for this alternate microcode are shown in Figure 18.10.

### 18.3 Multi-way Branches

As we saw in the previous section, using an instruction sequencer greatly reduces the size of our microcode memory but at the expense of restricting each state to have at most two next states (`upc+1` and `br_upc`). This restriction can be a problem if we have an FSM that has a large number of exits from a particular state. For example, in a microcoded processor, it is typical to branch to one of 10s to 100s of next states based on the *opcode* of the current instruction. Another multi-way branch is then needed on the *addressing mode* of the instruction. Implementing such a multi-way dispatch using the sequencer of Section 18.2 would result in very poor efficiency as  $n$  cycles would be required to test for  $n$  different opcodes.

We can overcome this limitation of two-way branches by using an instruction sequencer that supports multi-way branches as shown in Figure 18.11. This sequencer is similar to that of Figure 18.6 except that the branch target, `br_upc`, is generated from the branch instruction, `brinst`, and inputs rather than being provided directly from the microinstruction. With this approach, we can branch






Figure 18.11: A microcoded FSM with an instruction sequencer that supports multi-way branches.

Encoding	Opcode	Description
BRx	00xx	Branch on condition x. (As in Table 18.4 includes BR)
BRNx	01xx	Branch on not condition x. (As in Table 18.4 includes NOP)
BR4	1000	Branch four ways. $nupc = br\_upc + in$ .

Table 18.7: Branch instructions for microcoded FSM supporting multi-way branches.

to up to  $2^i$  next states (one for each input combination) from each state.

The branch instruction encodes not just the condition to test, but also how to determine the branch target. Table 18.7 shows one possible method of encoding multi-way branch instructions. The BRx and BRNx instructions are two-way branch instructions identical to those specified in Table 18.4. The BR4 instruction is a four-way branch that selects one of four adjacent states (from  $br\_upc$  to  $br\_upc+3$ ) depending on the input.

To use the BR4 instruction requires some care in mapping states to microinstruction addresses and may require that some states are duplicated. Consider, for example, the state diagram of Figure 18.12. A mapping of this state diagram to microcode addresses for a machine with a four-way branch instruction that adds the input to the branch target is shown in Table 18.8. The four-way branch from X targets address 000, so we must lay out branch targets A1, B1, C1, and X in locations 000, 001, 010, and 011 respectively. In a similar manner we locate states so that the four-way branch from C1 targets address 100. Thus we must place states C2, C3, X, and C1 at locations 100, 101, 110, and 111 respectively. To make this work we need two copies of X, one at 011 and one at




Figure 18.12: State diagram with two four-way branches.

Address	State	Branch Inst	Branch Target
000	A1	BR	A2
001	B1	BR	C2
010	C1	BR4	C2
011	X	BR4	A1
100	C2	BR	X
101	C3	BR	X
110	X'	BR4	A1
111	C1'	BR4	C2

Table 18.8: Mapping of state diagram of Figure 18.12 onto microcode addresses. States X and C1 are duplicated because they each appear in two four-way branches.

110 and two copies of C1 (at 010 and 111). When we duplicate a state in this manner we simply arrange for the two copies (e.g., X and X') to have identical behavior.

## 18.4 Multiple Instruction Types

So far we have considered microcoded FSMs that update all output bits in every microinstruction. In general, most FSMs need to update only a subset of the outputs in a given state. Our traffic-light controller FSMs, for example, change at most one light on each state change. We can save bits of the microinstruction by modifying our FSM to update only a single output register in any given state. We waste other microinstruction bits by specifying a branch instruction and branch target in each microinstruction even though many microinstructions




Figure 18.13: Instruction formats for a microcoded FSM with separate output and branch instructions.




Figure 18.14: Block diagram of a microcoded FSM with output instructions.

always proceed to the next state and don't branch. We can save these redundant branch bits by only branching in some instructions and updating outputs in other instructions.

Figure 18.13 shows the instruction formats for a microcoded FSM with two microinstruction types: a branch instruction and a store (output) instruction. Each microinstruction is one or the other. A branch instruction, identified by a 1 in the left-most bit, specifies a branch condition and a branch target. When the FSM encounters a branch microinstruction, it branches (or doesn't) as specified by the condition and target. No outputs are updated. A store instruction, identified by a 0 in the left-most bit, specifies an output register and a value. When the FSM encounters a store microinstruction it stores the value to the specified output register and then proceeds to the next microinstruction in sequence. No branching takes place with a store instruction.

Figure 18.14 shows a block diagram of a microcoded FSM that supports the two instruction types of Figure 18.13. Each microinstruction is split into an  $x$ -bit instruction field and a  $s$ -bit value field. The instruction field holds the operation-code (opcode) bit (the left-most bit that distinguishes between branch and store) and either the condition (for a branch) or the destination (for

a store). The value field holds either the branch target (for a branch) or the new output value (for a store).

The instruction sequencer of Figure 18.14 is identical to that of Figure 18.6 except that the branch logic always selects the next microinstruction `upc+1` when the current microinstruction is a store instruction.<sup>1</sup> The major difference is with the output logic. Here a decoder enables at most one output register to receive the value field on a store instruction.

A Verilog description for the microcode engine with two instruction types of Figure 18.14 is shown in Figure 18.15. Here the microinstruction is split into an opcode (0=store, 1=branch), an instruction (destination for store, condition for branch), and a value. For a store, the destination is decoded into a one-hot enable vector, `e`, that is used to enable the value to be stored into one of the three output registers (`ns=0`, `ew=1`, `lt=2`) or to load the timer (`dest = 3`). For a branch, `inst[2]` determines the polarity of the branch, and the low two bits, `inst[1:0]`, determine the condition to be tested (`lt=0`, `ew=1`, `lt|ew=2`, `timer=3`).

Table 18.9 shows the microcode for a more sophisticated traffic-light controller programmed on the microcode engine of Figure 18.15. The first three states load the three output registers with RED in the east-west and left-turn registers and GREEN in the north-south register. Next states NS1 and NS2 wait for eight cycles by loading the timer and waiting for it to signal done. State NS4 then waits for either input. The north-south light is set to YELLOW in NS5. NS6 and NS7 then set the timer and wait for it to be done before advancing to NS8 where the north-south light is set to RED. If the left-turn input is true NS9 branches to LT1 to sequence the left-turn light. Otherwise the east-west lights are sequenced in states EW1 through EW9. Waveforms from a simulation of this microcode are shown in Figure 18.16.

## 18.5 Microcode Subroutines

The state sequences in Table 18.9 are very repetitive. The NS, EW, and LT sequences perform largely the same actions. The only salient difference is the output register being written. Just as we shared common state sequences by factoring FSMs in Chapter 17, we can share common state sequences in a microcoded FSM by supporting *subroutines*. A subroutine is a sequence of instructions that can be called from several different points and after exiting returns control to the point from which it was called.

Figure 18.17 shows the block diagram of a microcode engine that supports one level of subroutines. This machine is identical to that of Figure 18.14 except for two differences: (a) a return uPC register, `rupc` and associated logic have been added to the sequencer, and (b) a select register and associated logic have been added to the output section.

---

<sup>1</sup>One can just as easily add multiple instruction types and output registers to a FSM that supports multi-way branches (Figure 18.11).



---

```

module ucodeMI(clk,rst,in,out) ;
    parameter n = 2 ; // input width
    parameter m = 9 ; // output width
    parameter o = 3 ; // output sub-width
    parameter k = 5 ; // bits of state
    parameter j = 4 ; // bits of instruction

    input  clk, rst ;
    input  [n-1:0] in ;
    output [m-1:0] out ;

    wire [k-1:0] nupc, upc ; // microprogram counter
    wire [j+k-1:0] uinst ; // microinstruction word
    wire done ; // timer done signal

    // split off fields of microinstruction
    wire opcode ; // opcode bit
    wire [j-2:0] inst ; // condition for branch, dest for store
    wire [k-1:0] value ; // target for branch, value for store
    assign {opcode, inst, value} = uinst ;

    DFF #(k) upc_reg(clk, nupc, upc) ; // microprogram counter
    ROM #(k,k+j) uc(upc, uinst) ; // microcode store

    // output registers and timer
    DFFE #(o) or0(clk, e[0], value[o-1:0], out[o-1:0]) ; // NS
    DFFE #(o) or1(clk, e[1], value[o-1:0], out[2*o-1:0]) ; // EW
    DFFE #(o) or2(clk, e[2], value[o-1:0], out[3*o-1:2*o]) ; // LT
    Timer #(k) tim(clk, rst, e[3], value, done) ; // timer

    // enable for output registers and timer
    wire [3:0] e = opcode ? 4'b0 : 1<<inst ;

    // branch instruction decode
    wire branch = opcode ? (inst[2] ^ (((inst[1:0] == 0) & in[0]) | // BLT
        ((inst[1:0] == 1) & in[1]) | // BEW
        ((inst[1:0] == 2) & (in[0]|in[1])) | //BLE
        ((inst[1:0] == 3) & done))) // BTD
        : 1'b0 ; // for a store opcode

    // microprogram counter
    assign nupc = rst ? {k{1'b0}} : branch ? value : upc + 1'b1 ;
endmodule

```

---

Figure 18.15: Verilog description of a microcoded FSM with two instruction types.

---

Address	State	Instruction	Value	Data
00000	RST1	SLT (0010)	RED 001	001000001
00001	RST2	SEW (0001)	RED 001	000100001
00010	NS1	SNS (0000)	GREEN 100	000000100
00011	NS2	STIM(0011)	TGRN 01000	001101000
00100	NS3	BNTD(1111)	NS3 00100	111100100
00101	NS4	BNLE(1110)	NS4 00101	111000101
00110	NS5	SNS (0000)	YELLOW 010	000000010
00111	NS6	STIM(0011)	TYEL 00011	001100011
01000	NS7	BNTD(1111)	NS7 01000	111101000
01001	NS8	SNS (0000)	RED 001	000000001
01010	NS9	BLT (1000)	LT1 10100	100010100
01011	EW1	STIM(0011)	TRED 00010	001100010
01100	EW2	BNTD(1111)	EW2 01100	111101100
01101	EW3	SEW (0001)	GREEN 100	000100100
01110	EW4	STIM(0011)	TGRN 01000	001101000
01111	EW5	BNTD(1111)	EW5 01111	111101111
10000	EW6	SEW (0001)	YELLOW 010	000100010
10001	EW7	STIM(0011)	TYEL 00011	001100011
10010	EW8	BNTD(1111)	EW8 10010	111110010
10011	EW9	BTD (1011)	RST2 00001	101100001
10100	LT1	STIM(0011)	TRED 00010	001100010
10101	LT2	BNTD(1111)	LT2 10101	111110101
10110	LT3	SLT (0010)	GREEN 100	001000100
10111	LT4	STIM(0011)	TGRN 01000	001101000
11000	LT5	BNTD(1111)	LT5 11000	111111000
11001	LT6	SLT (0010)	YELLOW 010	001000010
11010	LT7	STIM(0011)	TYEL 00011	001100011
11011	LT8	BNTD(1111)	LT8 10010	111111011
11100	LT9	BTD (1011)	RST1 00000	101100000

Table 18.9: Microcode to implement traffic-light controller on FSM with two instruction types.

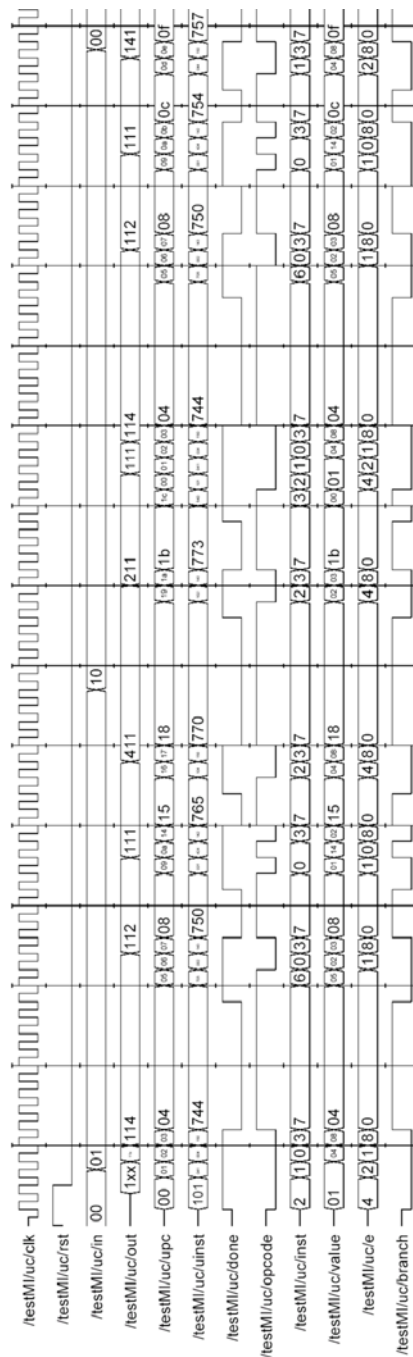


Figure 18.16: A waveform display showing a simulation of the microcoded FSM of Figure 18.15 using the microcode from Table 18.9.

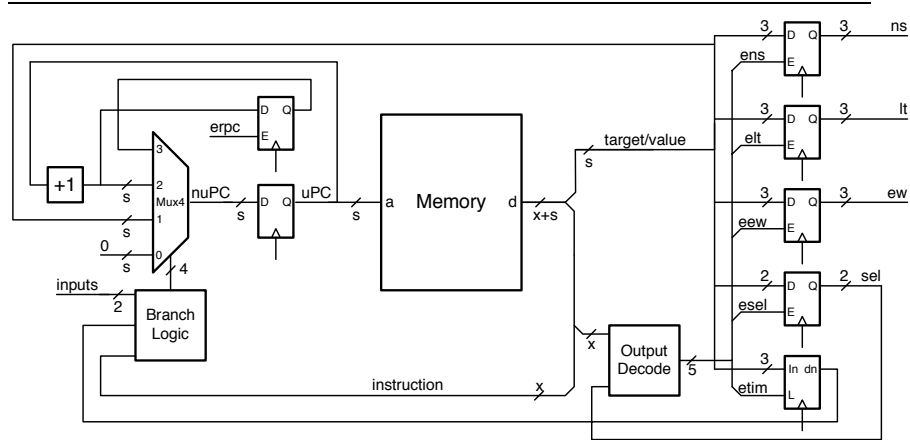


Figure 18.17: Microcoded FSM with support for one level of subroutines.

The **rupc** register is used to hold the **upc** to which a subroutine should branch when it completes. When a subroutine is called, the branch target is selected as the next **upc** and **upc+1**, the address of the next instruction in sequence is saved in the **rupc** register. A special branch instruction **CALL** is used to cause the enable line to the **rupc** register, **erpc**, to be asserted. When the subroutine is complete, it returns control to the saved location using another special branch instruction **RET** to select the **rupc** as the source of the next **upc**.

The select register is used to allow the same state sequence to write to different output registers when called from different places. A two-bit register identifier ( $ns=0$ ,  $ew=1$ ,  $lt=2$ ) can be stored in the select register. A special store instruction **SSEL** can then be used to store to the register specified by the select register (rather than by the destination bits of the instruction). Thus, the main program can store 0 (**ns**) into the select register and then call a subroutine to sequence the north-south lights on and off. The program can then store 1 (**ew**) into the select register and call the same subroutine to sequence the east-west lights on and off. The same subroutine can sequence different lights because it performs all of its output using the **SSEL** instruction.

EXAMPLE CODE HERE

## 18.6 A Simple Computer

## 18.7 Bibliographic Notes

## 18.8 Exercises

microcode sequence detector